

DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result, there are three main problems they need to solve:

- How to scale current manual processes and resources to screen 500,000 projects so that they can be posted as quickly and as efficiently as possible
- How to increase the consistency of project vetting across different volunteers to improve the experience for teachers
- How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature		Desc
<code>project_id</code>		A unique identifier for the proposed project. Example: p0
<code>project_title</code>	• •	Title of the project. Example: Art Will Make You Happy First Grade
<code>project_grade_category</code>	• • • •	Grade level of students for which the project is targeted. One of the following enumerated values: Preschool Kindergarten 1st Grade 2nd Grade 3rd Grade 4th Grade 5th Grade 6th Grade 7th Grade 8th Grade 9th Grade 10th Grade 11th Grade 12th Grade

Feature	Desc
	One or more (comma-separated) subject categories for the project from the following enumerated list of values:
project_subject_categories	<ul style="list-style-type: none"> Applied Learning Care & Health Health & Safety History & Culture Literacy & Language Math & Science Music & The Arts Special Education World Languages
	Example:
	<ul style="list-style-type: none"> Music & The Arts Literacy & Language, Math & Science
school_state	State where school is located (Two-letter U.S. postal code). (https://en.wikipedia.org/wiki/List_of_U.S._state_abbreviations#Postal_codes)
	Example:
project_subject_subcategories	One or more (comma-separated) subject subcategories for the project.
	Example:
	<ul style="list-style-type: none"> Literature & Writing Literature & Writing, Social Science
project_resource_summary	An explanation of the resources needed for the project. Example:
	<ul style="list-style-type: none"> My students need hands on literacy materials to meet sensory needs!<
project_essay_1	First application
project_essay_2	Second application
project_essay_3	Third application
project_essay_4	Fourth application
project_submitted_datetime	Datetime when project application was submitted. Example: 2016-01-12T12:43:50
teacher_id	A unique identifier for the teacher of the proposed project. Example: bdf8baa8fedef6bfeec7ae4ff1c
teacher_prefix	Teacher's title. One of the following enumerated values:
	<ul style="list-style-type: none"> Teacher Assistant Teacher Paraprofessional Volunteer Other
teacher_number_of_previously_posted_projects	Number of project applications previously submitted by the same teacher. Example:

* See the section **Notes on the Essay Data** for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25

Feature	Description
quantity	Quantity of the resource required. Example: 3
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
project_is_approved	A binary flag indicating whether DonorsChoose approved the project. A value of <code>0</code> indicates the project was not approved, and a value of <code>1</code> indicates the project was approved.

Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

- **project_essay_1:** "Introduce us to your classroom"
- **project_essay_2:** "Tell us more about your students"
- **project_essay_3:** "Describe how your students will use the materials you're requesting"
- **project_essay_3:** "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

- **project_essay_1:** "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."
- **project_essay_2:** "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

In [24]:

```
1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4
5 import sqlite3
6 import pandas as pd
7 import numpy as np
8 import nltk
9 import string
10 import matplotlib.pyplot as plt
11 import seaborn as sns
12 from sklearn.feature_extraction.text import TfidfTransformer
13 from sklearn.feature_extraction.text import TfidfVectorizer
14
15 from sklearn.feature_extraction.text import CountVectorizer
16 from sklearn.metrics import confusion_matrix
17 from sklearn import metrics
18 from sklearn.metrics import roc_curve, auc
19 from nltk.stem.porter import PorterStemmer
20
21 import re
22 # Tutorial about Python regular expressions: https://pymotw.com/2/re/
23 import string
24 from nltk.corpus import stopwords
25 from nltk.stem import PorterStemmer
26 from nltk.stem.wordnet import WordNetLemmatizer
27
28 from gensim.models import Word2Vec
29 from gensim.models import KeyedVectors
30 import pickle
31
32 from tqdm import tqdm
33 import os
34
35 from plotly import plotly
36 import plotly.offline as offline
37 import plotly.graph_objs as go
38 offline.init_notebook_mode()
39 from collections import Counter
```

1.1 Reading Data

In [25]:

```
1 project_data = pd.read_csv('train_data.csv')
2 resource_data = pd.read_csv('resources.csv')
```

In [26]:

```

1 print("Number of data points in train data", project_data.shape)
2 print('-'*50)
3 print("The attributes of data :", project_data.columns.values)

```

Number of data points in train data (109248, 17)

The attributes of data : ['Unnamed: 0' 'id' 'teacher_id' 'teacher_prefix' 'school_state'
'project_submitted_datetime' 'project_grade_category'
'project_subject_categories' 'project_subject_subcategories'
'project_title' 'project_essay_1' 'project_essay_2' 'project_essay_3'
'project_essay_4' 'project_resource_summary'
'teacher_number_of_previously_posted_projects' 'project_is_approved']

In [27]:

```

1 print("Number of data points in train data", resource_data.shape)
2 print(resource_data.columns.values)
3 resource_data.head(2)

```

Number of data points in train data (1541272, 4)

['id' 'description' 'quantity' 'price']

Out[27]:

	id	description	quantity	price
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	149.00
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	14.95

In [28]:

```

1
2 project_grade_category = []
3
4 for i in range(len(project_data)):
5     a = project_data["project_grade_category"][i].replace(" ", "_")
6     project_grade_category.append(a)
7
8

```

In [29]:

```
1 project_grade_category[0:5]
```

Out[29]:

['Grades_PreK-2', 'Grades_6-8', 'Grades_6-8', 'Grades_PreK-2', 'Grades_PreK-2']

In [30]:

```

1
2 project_data.drop(['project_grade_category'], axis=1, inplace=True)
3 project_data["project_grade_category"] = project_grade_category

```

1.1 Sorted by time

In [31]:

```

1 #https://stats.stackexchange.com/questions/341312/train-test-split-with-time-and-person
2 # how to replace elements in list python: https://stackoverflow.com/a/2582163/4084039
3 cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.co
4
5
6 #sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492/40840
7 project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
8 project_data.drop('project_submitted_datetime', axis=1, inplace=True)
9 project_data.sort_values(by=['Date'], inplace=True)
10
11
12 # how to reorder columns pandas python: https://stackoverflow.com/a/13148611/4084039
13 project_data = project_data[cols]
14
15
16 project_data.head(2)

```

Out[31]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	
55660	8393	p205479	2bf07ba08945e5d8b2a3f269b2b3cfe5	Mrs.	CA	00:
76127	37728	p043609	3f60494c61921b3b43ab61bdde2904df	Ms.	UT	00:

In []:

1

1.2 Adding resource data in dataframe

In [32]:

```

1 print("Number of data points in train data", resource_data.shape)
2 print(resource_data.columns.values)
3 resource_data.head(2)

```

Number of data points in train data (1541272, 4)
 ['id' 'description' 'quantity' 'price']

Out[32]:

	id	description	quantity	price
0	p233245	LC652 - Lakeshore Double-Space Mobile Drying Rack	1	149.00
1	p069063	Bouncy Bands for Desks (Blue support pipes)	3	14.95

In [33]:

```
1 price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_
2 project_data = pd.merge(project_data, price_data, on='id', how='left')
```

In [34]:

```
1 project_data.head(2)
```

Out[34]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	Date
0	8393	p205479	2bf07ba08945e5d8b2a3f269b2b3cfe5	Mrs.	CA	2016-04-27 00:27:36
1	37728	p043609	3f60494c61921b3b43ab61bdde2904df	Ms.	UT	2016-04-27 00:31:21

In [35]:

```
1 project_data = project_data.sample(n=50000)
2 #project_data=project_data.tail(1000)
3 project_data.shape
```

Out[35]:

(50000, 19)

In [36]:

```
1 project_data.shape
```

Out[36]:

(50000, 19)

1.2 preprocessing of project_subject_categories

In [37]:

```

1 categories = list(project_data['project_subject_categories'].values)
2 # remove special characters from list of strings python: https://stackoverflow.com/a/4
3
4 # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
5 # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-str
6 # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-pyt
7 cat_list = []
8 for i in categories:
9     temp = ""
10    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
11    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth
12        if 'The' in j.split(): # this will split each of the category based on space "
13            j=j.replace('The','') # if we have the words "The" we are going to replace
14            j = j.replace(' ', '') # we are placing all the ' '(space) with ''(empty) ex:"
15            temp+=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing sp
16            temp = temp.replace('&','_') # we are replacing the & value into
17    cat_list.append(temp.strip())
18
19 project_data['clean_categories'] = cat_list
20 project_data.drop(['project_subject_categories'], axis=1, inplace=True)
21
22 from collections import Counter
23 my_counter = Counter()
24 for word in project_data['clean_categories'].values:
25     my_counter.update(word.split())
26
27 cat_dict = dict(my_counter)
28 sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))
29

```

1.3 preprocessing of project_subject_subcategories

In [38]:

```

1 sub_categories = list(project_data['project_subject_subcategories'].values)
2 # remove special characters from list of strings python: https://stackoverflow.com/a/4
3
4 # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
5 # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-str
6 # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-pyt
7
8 sub_cat_list = []
9 for i in sub_categories:
10     temp = ""
11     # consider we have text like this "Math & Science, Warmth, Care & Hunger"
12     for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth
13         if 'The' in j.split(): # this will split each of the category based on space "
14             j=j.replace('The','') # if we have the words "The" we are going to replace
15             j = j.replace(' ', '') # we are placing all the ' '(space) with ''(empty) ex: "
16             temp +=j.strip()+" #" abc ".strip() will return "abc", remove the trailing sp
17             temp = temp.replace('&','_')
18     sub_cat_list.append(temp.strip())
19
20 project_data['clean_subcategories'] = sub_cat_list
21 project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)
22
23 # count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039
24 my_counter = Counter()
25 for word in project_data['clean_subcategories'].values:
26     my_counter.update(word.split())
27
28 sub_cat_dict = dict(my_counter)
29 sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

In [39]:

```
1 project_data.columns
```

Out[39]:

```

Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
      'Date', 'project_title', 'project_essay_1', 'project_essay_2',
      'project_essay_3', 'project_essay_4', 'project_resource_summary',
      'teacher_number_of_previously_posted_projects', 'project_is_approve
d',
      'project_grade_category', 'price', 'quantity', 'clean_categories',
      'clean_subcategories'],
      dtype='object')

```

In []:

```
1
```

1.3 Text preprocessing

In [40]:

```

1 # merge two column text dataframe:
2 project_data["essay"] = project_data["project_essay_1"].map(str) + \
3 project_data["project_essay_2"].map(str) + \
4 project_data["project_essay_3"].map(str) + \
5 project_data["project_essay_4"].map(str)

```

In [41]:

```
1 project_data.head(2)
```

Out[41]:

Unnamed: 0	id	teacher_id	teacher_prefix	school_state	
60490	163252	p069033	f7ac240660515b986030743e83db1766	Mrs.	LA
					13
42105	3358	p224685	e1df7b66db29045e36fbbc6f38a69b33	Ms.	CT
					19

In [42]:

```
1 ##### 1.4.2.3 Using Pretrained Models: TFIDF weighted W2V
```

In [43]:

```

1 '''# printing some random reviews
2 print(project_data['essay'].values[0])
3 print("="*50)
4 print(project_data['essay'].values[150])
5 print("="*50)
6 print(project_data['essay'].values[1000])
7 print("="*50)
8 print(project_data['essay'].values[20000])
9 print("="*50)
10 print(project_data['essay'].values[99999])
11 print("="*50)'''

```

Out[43]:

```

'# printing some random reviews\nprint(project_data[\'essay\'].values[0])\nprint("="*50)\nprint(project_data[\'essay\'].values[150])\nprint("="*50)\nprint(project_data[\'essay\'].values[1000])\nprint("="*50)\nprint(project_data[\'essay\'].values[20000])\nprint("="*50)\nprint(project_data[\'essay\'].values[99999])\nprint("="*50)'

```

In [44]:

```

1 # https://stackoverflow.com/a/47091490/4084039
2 import re
3
4 def decontracted(phrase):
5     # specific
6     phrase = re.sub(r"won't", "will not", phrase)
7     phrase = re.sub(r"can't", "can not", phrase)
8
9     # general
10    phrase = re.sub(r"n't", " not", phrase)
11    phrase = re.sub(r"'re", " are", phrase)
12    phrase = re.sub(r"'s", " is", phrase)
13    phrase = re.sub(r"'d", " would", phrase)
14    phrase = re.sub(r"'ll", " will", phrase)
15    phrase = re.sub(r"'t", " not", phrase)
16    phrase = re.sub(r"'ve", " have", phrase)
17    phrase = re.sub(r"'m", " am", phrase)
18    return phrase

```

In [45]:

```

1 ''sent = decontracted(project_data['essay'].values[20000])
2 print(sent)
3 print("="*50)'''

```

Out[45]:

```

'sent = decontracted(project_data['essay'].values[20000])\nprint(sent)\npr
int("="*50)'

```

In [47]:

```

1 '''# \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks
2 sent = sent.replace('\r', ' ')
3 sent = sent.replace('\\"', ' ')
4 sent = sent.replace('\n', ' ')
5 print(sent)'''

```

Out[47]:

```

'# \r \n \t remove from string python: http://texthandler.com/info/remove-li
ne-breaks-python/\nsent (http://texthandler.com/info/remove-line-breaks-pyth
on/\nsent) = sent.replace('\r', '\ ') \nsent = sent.replace('\\"', '\ '
') \nsent = sent.replace('\n', '\ ') \nprint(sent)'

```

In []:

```

1 '''#remove spacial character: https://stackoverflow.com/a/5843547/4084039
2 sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
3 print(sent)'''

```

```
1 # https://gist.github.com/sebleier/554280
2 # we are removing the words from the stop words list: 'no', 'nor', 'not'
3 stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're",
4             "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
5             'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'his',
6             'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll",
7             'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
8             'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until',
9             'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
10            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
11            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
12            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
13            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no',
14            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
15            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
16            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won',
17            'won't', 'wouldn', "wouldn't"]
```

In [50]:

```
1 # Combining all the above students
2 from tqdm import tqdm
3 preprocessed_essays = []
4 # tqdm is for printing the status bar
5 for sentence in tqdm(project_data['essay'].values):
6     sent = decontracted(sentence)
7     sent = sent.replace('\r', ' ')
8     sent = sent.replace('\n', ' ')
9     sent = sent.replace('\n', ' ')
10    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
11    # https://gist.github.com/sebleier/554280
12    sent = ' '.join(e for e in sent.split() if e not in stopwords)
13    preprocessed_essays.append(sent.lower().strip())
```

In [51]:

```
1 # after preprocessing
2 #preprocessed_essays[10:]
```

```
1 project_data['preprocessed_essays'] = preprocessed_essays
2 project_data.drop(['essay'], axis=1, inplace=True)
```

###

12/66

In [53]:

```
1 # similarly you can preprocess the titles also
```

In [54]:

```
1 # Combining all the above statemennts
2 from tqdm import tqdm
3 preprocessed_project_title = []
4 # tqdm is for printing the status bar
5 for sentence in tqdm(project_data['project_title'].values):
6     sent = decontracted(sentence)
7     sent = sent.replace('\\r', ' ')
8     sent = sent.replace('\\\"', ' ')
9     sent = sent.replace('\\n', ' ')
10    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
11    # https://gist.github.com/sebleier/554280
12    sent = ' '.join(e for e in sent.split() if e not in stopwords)
13    preprocessed_project_title.append(sent.lower().strip())
```

```
100%|████████████████████████████████████████████████████████████████████████████████|
50000/50000 [00:00<00:00, 51576.55it/s]
```

In [55]:

```
1 # after preprocessing
2 #preprocessed_project_title[1000]
```

In [56]:

```
1 #https://stackoverflow.com/questions/26666919/add-column-in-dataframe-from-list/384907
2 project_data['preprocessed_project_title'] = preprocessed_project_title
3 project_data.drop(['project_title'], axis=1, inplace=True)
```

In [57]:

```
1 project_data.head(2)
```

Out[57]:

	Unnamed: 0	id	teacher_id	teacher_prefix	school_state	
60490	163252	p069033	f7ac240660515b986030743e83db1766	Mrs.	LA	13
42105	3358	p224685	e1df7b66db29045e36fbbc6f38a69b33	Ms.	CT	19

In [58]:

```
1 project_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 50000 entries, 60490 to 106517
Data columns (total 20 columns):
Unnamed: 0                    50000 non-null int64
id                            50000 non-null object
teacher_id                   50000 non-null object
teacher_prefix               49998 non-null object
school_state                 50000 non-null object
Date                         50000 non-null datetime64[ns]
project_essay_1              50000 non-null object
project_essay_2              50000 non-null object
project_essay_3              1706 non-null object
project_essay_4              1706 non-null object
project_resource_summary     50000 non-null object
teacher_number_of_previously_posted_projects 50000 non-null int64
project_is_approved          50000 non-null int64
project_grade_category       50000 non-null object
price                        50000 non-null float64
quantity                     50000 non-null int64
clean_categories             50000 non-null object
clean_subcategories          50000 non-null object
preprocessed_essays          50000 non-null object
preprocessed_project_title    50000 non-null object
dtypes: datetime64[ns](1), float64(1), int64(4), object(14)
memory usage: 8.0+ MB
```

Apply KNN

1. [Task-1] Apply KNN(brute force version) on these feature sets


- **Set 1:** categorical, numerical features + project_title(BOW) + preprocessed_essay (BOW)
- **Set 2:** categorical, numerical features + project_title(TFIDF)+ preprocessed_essay (TFIDF)
- **Set 3:** categorical, numerical features + project_title(AVG W2V)+ preprocessed_essay (AVG W2V)
- **Set 4:** categorical, numerical features + project_title(TFIDF W2V)+ preprocessed_essay (TFIDF W2V)

2. Hyper paramter tuning to find best K


- Find the best hyper parameter which results in the maximum [AUC](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/>) value
- Find the best hyper paramter using k-fold cross validation (or) simple cross validation data
- Use gridsearch-cv or randomsearch-cv or write your own for loops to do this task

3. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, as shown in the figure



- Once you find the best hyper parameter, you need to train your model-M using the best hyper-param. Now, find the AUC on test data and plot the ROC curve on both train and test using model-M.



- Along with plotting ROC curve, you need to print the [confusion matrix](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>) with predicted and original labels of test data points

	Predicted: NO	Predicted: YES
Actual: NO	TN = ??	FP = ??
Actual: YES	FN = ??	TP = ??

4. [Task-2]

- Select top 2000 features from feature Set 2 using [SelectKBest](https://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.SelectKBest.html) (https://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.SelectKBest.html) and then apply KNN on top of these features

- ```

from sklearn.datasets import load_digits
from sklearn.feature_selection import SelectKBest, chi2
X, y = load_digits(return_X_y=True)
X.shape
X_new = SelectKBest(chi2, k=20).fit_transform(X, y)
X_new.shape
=====
output:
(1797, 64)
(1797, 20)

```

- Repeat the steps 2 and 3 on the data matrix after feature selection

## 5. Conclusion

- You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library [link \(http://zetcode.com/python/prettytable/\)](http://zetcode.com/python/prettytable/)

| Vectorizer | Model | Hyper parameter | AUC  |
|------------|-------|-----------------|------|
| BOW        | Brute | 7               | 0.78 |
| TFIDF      | Brute | 12              | 0.79 |
| W2V        | Brute | 10              | 0.78 |
| TFIDFW2V   | Brute | 6               | 0.78 |

### Note: Data Leakage

- There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
- To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
- While vectorizing your data, apply the method fit\_transform() on you train data, and apply the method transform() on cv/test data.
- For more details please go through this [link. \(https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf\)](https://soundcloud.com/applied-ai-course/leakage-bow-and-tfidf)

## 2. K Nearest Neighbor

### #2.1 Splitting data into Train and cross validation(or test): Stratified Sampling



In [59]:

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from sklearn import model_selection
5 from sklearn.model_selection import train_test_split
6 from sklearn.neighbors import KNeighborsClassifier
7 from sklearn.metrics import accuracy_score
8
9 from collections import Counter
10 from sklearn.metrics import accuracy_score
11
12 from sklearn.model_selection import cross_val_score
13 from sklearn.model_selection import cross_validate

```

In [60]:

```

1 y=project_data['project_is_approved']
2 y.shape

```

Out[60]:

(50000,)

In [61]:

```

1 #replace NAN to space https://stackoverflow.com/questions/49259305/raise-valueerrornp
2 project_data['teacher_prefix'] = project_data['teacher_prefix'].fillna(' ')

```

In [62]:

```

1 #https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html
2
3 #split the data into train and test for bag of words
4
5 x_train,x_test,y_train,y_test=model_selection.train_test_split(project_data,y,test_size=0.3,random_state=42)
6 #split train into cross val train and cross val test
7 x_train,x_cv,y_train,y_cv=model_selection.train_test_split(x_train,y_train,test_size=0.3,random_state=42)

```

splitting train\_data into train and cross validation in ratio of 7/3

In [63]:

```

1 # please write all the code with proper documentation, and proper titles for each sub-section
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpful in debugging your code
5 # when you plot any graph make sure you use
6 # a. Title, that describes your plot, this will be very helpful to the reader
7 # b. Legends if needed
8 # c. X-axis Label
9 # d. Y-axis Label

```

#2.2 Make Data Model Ready: encoding numerical, categorical features

## 2.2.1 encoding categorical features

In [64]:

```
1 x_train.head(2)
```

Out[64]:

|       | Unnamed:<br>0 | id      | teacher_id                       | teacher_prefix | school_state |    |
|-------|---------------|---------|----------------------------------|----------------|--------------|----|
| 39295 | 76565         | p155169 | aa5da503348fa16bc4c92658f4a28422 | Mr.            | MI           | 11 |
| 26606 | 102370        | p131104 | 6babe90a39e59f61ae6b9bd7297c41e4 | Mrs.           | PA           | 16 |

In [65]:

```
1 #one hot encoding for clean_categories
2 #
3 # we use count vectorizer to convert the values into one
4 from sklearn.feature_extraction.text import CountVectorizer
5 vectorizer = CountVectorizer(vocabulary=list(sorted_cat_dict.keys()), lowercase=False,
6 vectorizer.fit(x_train['clean_categories'].values)
7 x_train_categories_one_hot = vectorizer.transform(x_train['clean_categories'].values)
8 x_cv_categories_one_hot = vectorizer.transform(x_cv['clean_categories'].values)
9 x_test_categories_one_hot = vectorizer.transform(x_test['clean_categories'].values)
10 print(vectorizer.get_feature_names())
11 print("Shape of matrix after one hot encodig ",x_train_categories_one_hot.shape)
12 print("Shape of matrix after one hot encodig ",x_cv_categories_one_hot.shape)
13 print("Shape of matrix after one hot encodig ",x_test_categories_one_hot.shape)
14 print("*****50)
15
16 print(x_train_categories_one_hot[0:5])
```

```
['Warmth', 'Care_Hunger', 'History_Civics', 'Music_Arts', 'AppliedLearning',
'SpecialNeeds', 'Health_Sports', 'Math_Science', 'Literacy_Language']
```

```
Shape of matrix after one hot encodig (24500, 9)
```

```
Shape of matrix after one hot encodig (10500, 9)
```

```
Shape of matrix after one hot encodig (15000, 9)
```

```

```

```
(0, 5) 1
(0, 6) 1
(1, 6) 1
(2, 2) 1
(2, 7) 1
(3, 7) 1
(4, 7) 1
```

In [66]:

```

1 #one hot encoding for clean_subcategories
2 #
3 # we use count vectorizer to convert the values into one
4 from sklearn.feature_extraction.text import CountVectorizer
5 vectorizer = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=False)
6 vectorizer.fit(x_train['clean_subcategories'].values)
7 x_train_subcategories_one_hot = vectorizer.transform(x_train['clean_subcategories'].values)
8 x_cv_subcategories_one_hot = vectorizer.transform(x_cv['clean_subcategories'].values)
9 x_test_subcategories_one_hot = vectorizer.transform(x_test['clean_subcategories'].values)
10 print(vectorizer.get_feature_names())
11 print("Shape of matrix after one hot encoding ",x_train_subcategories_one_hot.shape)
12 print("Shape of matrix after one hot encoding ",x_cv_subcategories_one_hot.shape)
13 print("Shape of matrix after one hot encoding ",x_test_subcategories_one_hot.shape)

```

```

['Economics', 'CommunityService', 'FinancialLiteracy', 'ParentInvolvement',
 'Civics_Government', 'Extracurricular', 'ForeignLanguages', 'Warmth', 'Care_
Hunger', 'NutritionEducation', 'PerformingArts', 'SocialSciences', 'Characte
rEducation', 'TeamSports', 'Other', 'College_CareerPrep', 'Music', 'History_
Geography', 'Health_LifeScience', 'EarlyDevelopment', 'ESL', 'Gym_Fitness',
 'EnvironmentalScience', 'VisualArts', 'Health_Wellness', 'AppliedSciences',
 'SpecialNeeds', 'Literature_Writing', 'Mathematics', 'Literacy']

```

Shape of matrix after one hot encoding (24500, 30)

Shape of matrix after one hot encoding (10500, 30)

Shape of matrix after one hot encoding (15000, 30)

In [67]:

```

1 #one hot encoding for school_state
2
3
4 my_counter = Counter()
5 for state in project_data['school_state'].values:
6 my_counter.update(state.split())
7
8 school_state_cat_dict = dict(my_counter)
9 sorted_school_state_cat_dict = dict(sorted(school_state_cat_dict.items(), key=lambda k:
10
11 #
12 # we use count vectorizer to convert the values into one
13 from sklearn.feature_extraction.text import CountVectorizer
14 vectorizer = CountVectorizer(vocabulary=list(sorted_school_state_cat_dict.keys()), low
15 vectorizer.fit(x_train['school_state'].values)
16 x_train_school_state_one_hot = vectorizer.transform(x_train['school_state'].values)
17 x_cv_school_state_one_hot = vectorizer.transform(x_cv['school_state'].values)
18 x_test_school_state_one_hot = vectorizer.transform(x_test['school_state'].values)
19 print(vectorizer.get_feature_names())
20 print("Shape of matrix after one hot encoding ",x_train_school_state_one_hot.shape)
21 print("Shape of matrix after one hot encoding ",x_cv_school_state_one_hot.shape)
22 print("Shape of matrix after one hot encoding ",x_test_school_state_one_hot.shape)

```

```

['WY', 'VT', 'ND', 'MT', 'AK', 'RI', 'SD', 'NE', 'DE', 'NH', 'ME', 'HI', 'D
C', 'WV', 'NM', 'IA', 'KS', 'ID', 'AR', 'CO', 'MN', 'OR', 'NV', 'KY', 'MS',
 'MD', 'CT', 'TN', 'UT', 'AL', 'WI', 'VA', 'AZ', 'NJ', 'OK', 'LA', 'MA', 'W
A', 'OH', 'MO', 'IN', 'MI', 'PA', 'GA', 'SC', 'IL', 'NC', 'FL', 'NY', 'TX',
 'CA']

```

Shape of matrix after one hot encoding (24500, 51)

Shape of matrix after one hot encoding (10500, 51)

Shape of matrix after one hot encoding (15000, 51)

In [68]:

```

1 #one hot encoding for project_grade_category
2
3 my_counter = Counter()
4 for project_grade in project_data['project_grade_category'].values:
5 my_counter.update(project_grade.split())
6
7
8 project_grade_cat_dict = dict(my_counter)
9 sorted_project_grade_cat_dict = dict(sorted(project_grade_cat_dict.items(), key=lambda
10 #
11 # we use count vectorizer to convert the values into one
12 from sklearn.feature_extraction.text import CountVectorizer
13 vectorizer = CountVectorizer(vocabulary=list(sorted_project_grade_cat_dict.keys()), low
14 vectorizer.fit(x_train['project_grade_category'].values)
15 x_train_grade_category_one_hot = vectorizer.transform(x_train['project_grade_category'
16 x_cv_grade_category_one_hot = vectorizer.transform(x_cv['project_grade_category'].valu
17 x_test_grade_category_one_hot = vectorizer.transform(x_test['project_grade_category'].
18 print(vectorizer.get_feature_names())
19 print("Shape of matrix after one hot encodig ",x_train_grade_category_one_hot.shape)
20 print("Shape of matrix after one hot encodig ",x_cv_grade_category_one_hot.shape)
21 print("Shape of matrix after one hot encodig ",x_test_grade_category_one_hot.shape)

```

```
['Grades_9-12', 'Grades_6-8', 'Grades_3-5', 'Grades_PreK-2']
```

```
Shape of matrix after one hot encodig (24500, 4)
```

```
Shape of matrix after one hot encodig (10500, 4)
```

```
Shape of matrix after one hot encodig (15000, 4)
```

In [69]:

```

1 #one hot encoding for prefix_category
2
3 my_counter = Counter()
4 for teacher_prefix in project_data['teacher_prefix'].values:
5 teacher_prefix = str(teacher_prefix)
6 my_counter.update(teacher_prefix.split())
7
8 teacher_prefix_cat_dict = dict(my_counter)
9 sorted_teacher_prefix_cat_dict = dict(sorted(teacher_prefix_cat_dict.items(), key=lambda
10 #
11 # we use count vectorizer to convert the values into one
12 from sklearn.feature_extraction.text import CountVectorizer
13 vectorizer = CountVectorizer(vocabulary=list(sorted_teacher_prefix_cat_dict.keys()), low
14 vectorizer.fit(x_train['teacher_prefix'].values)
15 x_train_prefix_one_hot = vectorizer.transform(x_train['teacher_prefix'].values)
16 x_cv_prefix_one_hot = vectorizer.transform(x_cv['teacher_prefix'].values)
17 x_test_prefix_one_hot = vectorizer.transform(x_test['teacher_prefix'].values)
18 print(vectorizer.get_feature_names())
19 print("Shape of matrix after one hot encodig ",x_train_prefix_one_hot.shape)
20 print("Shape of matrix after one hot encodig ",x_cv_prefix_one_hot.shape)
21 print("Shape of matrix after one hot encodig ",x_test_prefix_one_hot.shape)
22
23
24

```

```
['Dr.', 'Teacher', 'Mr.', 'Ms.', 'Mrs.']
```

```
Shape of matrix after one hot encodig (24500, 5)
```

```
Shape of matrix after one hot encodig (10500, 5)
```

```
Shape of matrix after one hot encodig (15000, 5)
```

In [70]:

```

1 # please write all the code with proper documentation, and proper titles for each subse
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpfull in debugging you
5 # make sure you featurize train and test data separatly
6
7 # when you plot any graph make sure you use
8 # a. Title, that describes your plot, this will be very helpful to the reader
9 # b. Legends if needed
10 # c. X-axis Label
11 # d. Y-axis Label

```

## 2.2.2 encoding numerical features

In [71]:

```
1 x_train.head(2)
```

Out[71]:

|       | Unnamed: 0 | id      | teacher_id                       | teacher_prefix | school_state |    |
|-------|------------|---------|----------------------------------|----------------|--------------|----|
| 39295 | 76565      | p155169 | aa5da503348fa16bc4c92658f4a28422 | Mr.            | MI           | 11 |
| 26606 | 102370     | p131104 | 6babe90a39e59f61ae6b9bd7297c41e4 | Mrs.           | PA           | 16 |

In [72]:

```

1 #price standardization of x_train data
2 #-----
3 # check this one: https://www.youtube.com/watch?v=0H0q0cLn3Z4&t=530s
4 # standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler
5 from sklearn.preprocessing import StandardScaler
6
7 # price_standardized = standardScaler.fit(project_data['price'].values)
8 # this will rise the error
9 # ValueError: Expected 2D array, got 1D array instead: array=[725.05 213.03 329. ...]
10 # Reshape your data either using array.reshape(-1, 1)
11
12 price_scalar = StandardScaler()
13 price_scalar.fit(x_train['price'].values.reshape(-1,1)) # finding the mean and standard deviation
14 x_train_price_standardized=price_scalar.transform(x_train['price'].values.reshape(-1,1))
15 x_cv_price_standardized = price_scalar.transform(x_cv['price'].values.reshape(-1, 1))
16 x_test_price_standardized = price_scalar.transform(x_test['price'].values.reshape(-1, 1))
17
18 print(f"Mean : {price_scalar.mean_[0]}, Standard deviation : {np.sqrt(price_scalar.var_[0])}")
19
20 print("After standardization")
21 print(x_train_price_standardized.shape, y_train.shape)
22 print(x_cv_price_standardized.shape, y_cv.shape)
23 print(x_test_price_standardized.shape, y_test.shape)
24
25 # Now standardize the data with above mean and variance.
26 #x_train_price_standardized = price_scalar.transform(x_train['price'].values.reshape(-1, 1))

```

Mean : 297.4391379591837, Standard deviation : 355.84486777404317

After standardization

(24500, 1) (24500,)

(10500, 1) (10500,)

(15000, 1) (15000,)

## 2.2.3 merge numerical and categorical data

In [73]:

```

1 # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
2 from scipy.sparse import hstack
3 # with the same hstack function we are concatenating a sparse matrix and a dense matrix
4 x_train_ohe = hstack((x_train_categories_one_hot, x_train_subcategories_one_hot, x_train_price_standardized))
5 x_cv_ohe = hstack((x_cv_categories_one_hot, x_cv_subcategories_one_hot, x_cv_price_standardized))
6 x_test_ohe = hstack((x_test_categories_one_hot, x_test_subcategories_one_hot, x_test_price_standardized))
7
8 print(x_train_ohe.shape)
9 print(x_cv_ohe.shape)
10 print(x_test_ohe.shape)

```

(24500, 100)

(10500, 100)

(15000, 100)

In [74]:

```
1 type(x_train_ohe)
```

Out[74]:

```
scipy.sparse.csr.csr_matrix
```

## 2.3 Make Data Model Ready: encoding eassay, and project\_title

## 2.4 Applying KNN on different kind of featurization as mentioned in the instructions

In [75]:

```
1 # please write all the code with proper documentation, and proper titles for each subse
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpfull in debugging your
5
6 # when you plot any graph make sure you use
7 # a. Title, that describes your plot, this will be very helpful to the reader
8 # b. Legends if needed
9 # c. X-axis Label
10 # d. Y-axis Label
```

### 2.4.1 Applying KNN brute force on BOW, SET 1

#### vectorize the essay and title data, SET 1

In [76]:

```
1 #you can vectorize the essay
2 #
3 #https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.Count
4
5 # We are considering only the words which appeared in at least 10 documents(rows or pro
6 vectorizer = CountVectorizer(min_df=10, ngram_range=(1,4), max_features=5000)
7 vectorizer.fit(x_train['preprocessed_essays'].values)# fit has to apply only on train c
8
9 # we use fitted CountVectorizer to convert the text to vector
10 x_train_bow_essays = vectorizer.transform(x_train['preprocessed_essays'].values)
11 x_cv_bow_essays = vectorizer.transform(x_cv['preprocessed_essays'].values)
12 x_test_bow_essays = vectorizer.transform(x_test['preprocessed_essays'].values)
13
14 print("Shape of matrix after one hot encodig ",x_train_bow_essays.shape, y_train.shape)
15 print("Shape of matrix after one hot encodig ",x_cv_bow_essays.shape)
16 print("Shape of matrix after one hot encodig ",x_test_bow_essays.shape)
```

```
Shape of matrix after one hot encodig (24500, 5000) (24500,)
```

```
Shape of matrix after one hot encodig (10500, 5000)
```

```
Shape of matrix after one hot encodig (15000, 5000)
```

In [77]:

```

1 #https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVec
2 #you can vectorize the title
3 # We are considering only the words which appeared in at least 10 documents(rows or pro
4 vectorizer = CountVectorizer(min_df=10, ngram_range=(1,4), max_features=5000)
5 vectorizer.fit(x_train['preprocessed_project_title'].values)# fit has to apply only on
6
7 # we use fitted CountVectorizer to convert the text to vector
8 x_train_bow_title = vectorizer.transform(x_train['preprocessed_project_title'].values)
9 x_cv_bow_title = vectorizer.transform(x_cv['preprocessed_project_title'].values)
10 x_test_bow_title = vectorizer.transform(x_test['preprocessed_project_title'].values)
11
12 print("Shape of matrix after one hot encodig ",x_train_bow_title.shape)
13 print("Shape of matrix after one hot encodig ",x_cv_bow_title.shape)
14 print("Shape of matrix after one hot encodig ",x_test_bow_title.shape)

```

Shape of matrix after one hot encodig (24500, 2142)

Shape of matrix after one hot encodig (10500, 2142)

Shape of matrix after one hot encodig (15000, 2142)

In [78]:

```
1 x_train.columns
```

Out[78]:

```

Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
 'Date', 'project_essay_1', 'project_essay_2', 'project_essay_3',
 'project_essay_4', 'project_resource_summary',
 'teacher_number_of_previously_posted_projects', 'project_is_approve
d',
 'project_grade_category', 'price', 'quantity', 'clean_categories',
 'clean_subcategories', 'preprocessed_essays',
 'preprocessed_project_title'],
 dtype='object')

```

In [79]:

```
1 # Please write all the code with proper documentation
```

**merge dataset, SET 1**



In [80]:

```

1 # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
2 #https://stackoverflow.com/questions/30163830/accessing-elements-in-coo-matrix
3 from scipy.sparse import hstack
4 from scipy.sparse import coo_matrix
5 # with the same hstack function we are concatenating a sparse matrix and a dense matrix
6 x_train_bow = hstack((x_train_ohe, x_train_bow_essays, x_train_bow_title)).tocsr()
7 x_cv_bow = hstack((x_cv_ohe, x_cv_bow_essays, x_cv_bow_title)).tocsr()
8 x_test_bow = hstack((x_test_ohe, x_test_bow_essays, x_test_bow_title)).tocsr()
9
10 print(x_train_bow.shape, y_train.shape)
11 print(x_cv_bow.shape)
12 print(x_test_bow.shape)

```

(24500, 7242) (24500,)

(10500, 7242)

(15000, 7242)

In [ ]:

```
1 type(x_train_bow)
```

## Hyperparameter tuning by AUC plot for cv and train dataset, SET 1


In [81]:

```

1 def batch_predict(clf, data):
2 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
3 # not the predicted outputs
4
5 y_data_pred = []
6 tr_loop = data.shape[0] - data.shape[0]%1000
7 # consider you X_tr shape is 49041, then your tr_loop will be 49041 - 49041%1000 =
8 # in this for loop we will iterate until the last 1000 multiplier
9 for i in range(0, tr_loop, 1000):
10 y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
11 # we will be predicting for the last data points
12 if data.shape[0]%1000 != 0:
13 y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])
14
15 return y_data_pred

```





## PARAMETER TUNING USING GRID SEARCH

In [ ]:

```

1 '''# https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
2 from sklearn.model_selection import GridSearchCV
3
4 neigh = KNeighborsClassifier()
5
6 grid_val = {'n_neighbors':[1, 5, 10, 15, 20, 30, 40, 50, 60, 70, 80, 90, 99]}
7
8 clf = GridSearchCV(neigh, grid_val, cv= 5, scoring='roc_auc')
9 clf.fit(x_train_bow, y_train)
10
11 results_grid_bow = pd.DataFrame.from_dict(clf.cv_results_).sort_values(['param_n_neighbors'])
12
13
14 train_auc= clf.cv_results_['mean_train_score']
15 train_auc_std= clf.cv_results_['std_train_score']
16 cv_auc = clf.cv_results_['mean_test_score']
17 cv_auc_std= clf.cv_results_['std_test_score']
18
19 plt.plot(grid_val['n_neighbors'], train_auc, label='Train AUC')
20 # code reference: https://stackoverflow.com/a/48803361/4084039
21 plt.gca().fill_between(grid_val['n_neighbors'],train_auc - train_auc_std,train_auc + train_auc_std)
22
23 plt.plot(grid_val['n_neighbors'], cv_auc, label='CV AUC')
24 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
25 plt.gca().fill_between(grid_val['n_neighbors'],cv_auc - cv_auc_std,cv_auc + cv_auc_std)
26
27 plt.scatter(grid_val['n_neighbors'], train_auc, label='Train AUC points')
28 plt.scatter(grid_val['n_neighbors'], cv_auc, label='CV AUC points')
29
30 plt.legend()
31 plt.xlabel("K: hyperparameter")
32 plt.ylabel("AUC")
33 plt.title("ERROR PLOTS")
34 plt.grid()
35 plt.show()
36
37 results_grid_bow.head()'''

```

In [ ]:

1

## PARAMETER TUNING USING Random search

In [ ]:

```

1 '''# https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
2 from sklearn.model_selection import GridSearchCV
3 from scipy.stats import randint as sp_randint
4 from sklearn.model_selection import RandomizedSearchCV
5
6 neigh = KNeighborsClassifier(n_jobs=-1)
7 parameters = {'n_neighbors':sp_randint(40, 100)}
8 clf = RandomizedSearchCV(neigh, parameters, cv=3, scoring='roc_auc')
9 clf.fit(x_train_bow, y_train)
10
11 results_rand_bow = pd.DataFrame.from_dict(clf.cv_results_).sort_values(['param_n_neighbors'])
12
13
14 train_auc= results['mean_train_score']
15 train_auc_std= results['std_train_score']
16 cv_auc = results['mean_test_score']
17 cv_auc_std= results['std_test_score']
18 K = results['param_n_neighbors']
19
20 plt.plot(K, train_auc, label='Train AUC')
21 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
22 # plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_std,alpha=0.2,color='d')
23
24 plt.plot(K, cv_auc, label='CV AUC')
25 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
26 # plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2,color='d')
27
28 plt.scatter(K, train_auc, label='Train AUC points')
29 plt.scatter(K, cv_auc, label='CV AUC points')
30
31
32 plt.legend()
33 plt.xlabel("K: hyperparameter")
34 plt.ylabel("AUC")
35 plt.title("Hyper parameter Vs AUC plot")
36 plt.grid()
37 plt.show()
38
39 results_rand_bow.head()'''

```

In [128]:

```

1 # from the error plot we choose K such that, we will have maximum AUC on cv data and go
2 # Note: based on the method you use you might get different hyperparameter values as be
3 # so, you choose according to the method you choose, you use gridsearch if you are hav
4 # if you increase the cv values in the GridSearchCV you will get more robust results.
5
6 #here we are choosing the best_k based on forLoop results
7 opt_k_bow=90

```

## Apply best hyperparameter on test dataset, SET 1

You need to plot the performance of model both on train data and cross validation data for each hyperparameter, as shown in the figure

Once you find the best hyper parameter, you need to train your model-M using the best hyper-param. Now, find


the AUC on test data and plot the ROC curve on both train and test using model-M.

Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels of test data points

IF Your system is getting stuck when you are working with the Knn: YOU NEED TO USE BATCH WISE PREDICTION

In [129]:

```
1 # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
2 from sklearn.metrics import roc_curve, auc
3
4
5 neigh = KNeighborsClassifier(n_neighbors=opt_k_bow, n_jobs=-1)
6 neigh.fit(x_train_bow, y_train)
7 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
8 # not the predicted outputs
9
10 y_train_pred = batch_predict(neigh, x_train_bow)
11 y_test_pred = batch_predict(neigh, x_test_bow)
12
13 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
14 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
15 tain_auc_bow=auc(train_fpr, train_tpr)
16 test_auc_bow=auc(test_fpr, test_tpr)
17
18 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
19 plt.plot(test_fpr, test_tpr, label="train AUC =" +str(auc(test_fpr, test_tpr)))
20 plt.legend()
21 plt.xlabel("K: hyperparameter")
22 plt.ylabel("AUC")
23 plt.title("ERROR PLOTS")
24 plt.grid()
25 plt.show()
```



In [ ]:

1

In [ ]:

```

1 '''knn1 = KNeighborsClassifier(n_neighbors=opt_k_bow,algorithm='brute',weights='uniform')
2 knn1.fit(x_train_bow,y_train)
3
4 pred_prob_test = knn1.predict_proba(x_test_bow)
5
6
7 #AUC of train dataset
8 pred_prob_train = knn1.predict_proba(x_train_bow)
9 fpr1, tpr1, thresholds = roc_curve(y_train,pred_prob_train[:, 1])
10 bow_roc_auc_train = auc(fpr1, tpr1)
11 print("Best AUC of train: ",bow_roc_auc_train)
12
13 #AUC of test dataset
14 pred_prob_test = knn1.predict_proba(x_test_bow)
15 fpr2, tpr2, thresholds = roc_curve(y_test,pred_prob_test[:, 1])
16 bow_roc_auc_test = auc(fpr2, tpr2)
17 print("Best AUC of test: ",bow_roc_auc_test)
18
19 #value taken from from GridsearchCV section
20 plt.title('Receiver Operating Characteristic')
21 plt.plot(fpr1, tpr1, 'r',label='AUC_train = %.2f'% bow_roc_auc_train)
22 plt.plot(fpr2, tpr2, 'g',label='AUC_test = %.2f'% bow_roc_auc_test)
23 plt.legend(loc='lower right')
24 plt.plot([0,1],[0,1],'r--')
25 plt.xlim([-0.1,1.1])
26 plt.ylim([-0.1,1.1])
27 plt.ylabel('True Positive Rate')
28 plt.xlabel('False Positive Rate')
29 plt.show()
30
31 print("Best AUC of train: ",bow_roc_auc_train)
32 print("Best AUC of test: ",bow_roc_auc_test)'''

```

In [130]:

```

1 # we are writing our own function for predict, with defined threshold
2 # we will pick a threshold that will give the least fpr
3 def find_best_threshold(threshold, fpr, tpr):
4 t = threshold[np.argmax(tpr*(1-fpr))]
5 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
6 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.rou
7 return t
8
9 def predict_with_best_t(proba, threshold):
10 predictions = []
11 for i in proba:
12 if i>=threshold:
13 predictions.append(1)
14 else:
15 predictions.append(0)
16 return predictions

```

In [131]:

```

1 print("="*100)
2 from sklearn.metrics import confusion_matrix
3 best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
4 print("Train confusion matrix")
5 print(confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t)))
6 print("Test confusion matrix")
7 print(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)))

```

```

=====
=====
the maximum value of tpr*(1-fpr) 0.3844324967040373 for threshold 0.811
Train confusion matrix
[[2309 1393]
 [7979 12819]]
Test confusion matrix
[[1275 1025]
 [5033 7667]]

```

In [132]:

```

1 #CONFUSION MATRIX
2 def predict(proba, threshold, fpr, tpr):
3
4 t = threshold[np.argmax(fpr*(1-tpr))]
5
6 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
7
8 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.rou
9 predictions = []
10 for i in proba:
11 if i>=t:
12 predictions.append(1)
13 else:
14 predictions.append(0)
15 return predictions
16
17
18
19

```



In [133]:

```

1 print("="*100)
2 print("TRAIN confusion matrix")
3 print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, train_
4
5
6 conf_matr_df_train_1 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pred, tr
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```

TRAIN confusion matrix

the maximum value of  $tpr \cdot (1 - fpr)$  0.3844324967040373 for threshold 0.811


```
[[2309 1393]
```

```
 [7979 12819]]
```

the maximum value of  $tpr \cdot (1 - fpr)$  0.3844324967040373 for threshold 0.811

Out[133]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea2ad41860>



In [134]:

```

1
2 print("="*100)
3 print("Test confusion matrix")
4 print(confusion_matrix(y_test, predict(y_test_pred, te_thresholds, test_fpr, test_tpr))
5
6
7 conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred, te_th
8
9 sns.set(font_scale=1.4)#for label size
10 sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g')

```


```

=====
=====
Test confusion matrix
the maximum value of tpr*(1-fpr) 0.33466021910304694 for threshold 0.822
[[1462 838]
 [6190 6510]]
the maximum value of tpr*(1-fpr) 0.33466021910304694 for threshold 0.822

```

Out[134]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea20f59c18>



In [ ]:

1

## Feature selection with SelectKBest (top 2000) , SET 1

In [135]:

```
1 #https://scikit-learn.org/stable/modules/generated/sklearn.feature_selection.f_classif
2 #https://stackoverflow.com/questions/49300193/feature-selection-f-classif-scikit-learn
3 from sklearn.feature_selection import SelectKBest, chi2
4 from sklearn.feature_selection import f_classif
5
6 x_train_bow_2000 = SelectKBest(f_classif, k=2000).fit_transform(x_train_bow, y_train)
7 x_cv_bow_2000 = SelectKBest(f_classif, k=2000).fit_transform(x_cv_bow, y_cv)
8 x_test_bow_2000 = SelectKBest(f_classif, k=2000).fit_transform(x_test_bow, y_test)
9 print(x_train_bow_2000.shape)
10 print(x_cv_bow_2000.shape)
11 print(x_test_bow_2000.shape)
12
13
```

```
(24500, 2000)
(10500, 2000)
(15000, 2000)
```

In [136]:

```
1 type(x_train_bow_2000)
```

Out[136]:

```
scipy.sparse.csr.csr_matrix
```

## Parameter tuning using Gridsearch


In [137]:

```

1 train_auc = []
2 cv_auc = []
3 K = [1, 5, 10, 15, 20, 30, 40, 50, 60, 70, 80, 90, 99]
4 for i in tqdm(K):
5 neigh = KNeighborsClassifier(n_neighbors=i, n_jobs=-1)
6 neigh.fit(x_train_bow_2000, y_train)
7
8
9
10 y_train_pred = batch_predict(neigh, x_train_bow_2000)
11 y_cv_pred = batch_predict(neigh, x_cv_bow_2000)
12
13 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
14 # not the predicted outputs
15 train_auc.append(roc_auc_score(y_train, y_train_pred))
16 cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
17
18 plt.plot(K, train_auc, label='Train AUC')
19 plt.plot(K, cv_auc, label='CV AUC')
20
21 plt.scatter(K, train_auc, label='Train AUC points')
22 plt.scatter(K, cv_auc, label='CV AUC points')
23
24 plt.legend()
25 plt.xlabel("K: hyperparameter")
26 plt.ylabel("AUC")
27 plt.title("ERROR PLOTS")
28 plt.grid()
29 plt.show()

```

```
100%|███████████|
██████████ | 13/13 [54:14<00:00, 256.43s/it]
```



In [138]:

```
1 opt k bow 2000=40
```


## Applying on Test data and ROC

In [139]:

```

1 # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#skle
2 from sklearn.metrics import roc_curve, auc
3
4
5 neigh = KNeighborsClassifier(n_neighbors=opt_k_bow_2000, n_jobs=-1)
6 neigh.fit(x_train_bow_2000, y_train)
7 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of t
8 # not the predicted outputs
9
10 y_train_pred = batch_predict(neigh, x_train_bow_2000)
11 y_test_pred = batch_predict(neigh, x_test_bow_2000)
12
13 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
14 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
15 tain_auc_bow_2000=auc(train_fpr, train_tpr)
16 test_auc_bow_2000=auc(test_fpr, test_tpr)
17
18 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
19 plt.plot(test_fpr, test_tpr, label="train AUC =" +str(auc(test_fpr, test_tpr)))
20 plt.plot([0,1],[0,1], 'r--')
21 plt.legend()
22 plt.xlabel("K: hyperparameter")
23 plt.ylabel("AUC")
24 plt.title("ERROR PLOTS")
25 plt.grid()
26 plt.show()

```



In [140]:

```

1 #CONFUSION MATRIX
2 def predict(proba, threshold, fpr, tpr):
3
4 t = threshold[np.argmax(fpr*(1-tpr))]
5
6 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
7
8 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.ro
9 predictions = []
10 for i in proba:
11 if i>=t:
12 predictions.append(1)
13 else:
14 predictions.append(0)
15 return predictions
16
17
18

```

In [141]:

```

1 print("="*100)
2 print("TRAIN confusion matrix")
3 print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, train_
4
5
6 conf_matr_df_train_1 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pred, tr
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='g')

```


```

=====
=====
TRAIN confusion matrix
the maximum value of tpr*(1-fpr) 0.41908262799445295 for threshold 0.775
[[2438 1264]
 [7563 13235]]
the maximum value of tpr*(1-fpr) 0.41908262799445295 for threshold 0.775

```

Out[141]:

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x1ea2ad39ba8&gt;



In [142]:

```

1
2 print("="*100)
3 print("Test confusion matrix")
4 print(confusion_matrix(y_test, predict(y_test_pred, te_thresholds, test_fpr, test_tpr))
5
6
7 conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred, te_th
8
9 sns.set(font_scale=1.4)#for label size
10 sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```

Test confusion matrix

the maximum value of  $tpr \cdot (1 - fpr)$  0.2619001027045532 for threshold 0.675


```
[[1257 1043]
```

```
[6614 6086]]
```

the maximum value of  $tpr \cdot (1 - fpr)$  0.2619001027045532 for threshold 0.675

Out[142]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea21a7fc18>



## 2.4.2 Applying KNN brute force on TFIDF, SET 2

### TFIDF Vectorizing essay and title variable, SET 2

In [ ]:

```
1 # Please write all the code with proper documentation
```

In [83]:

```

1 from sklearn.feature_extraction.text import TfidfVectorizer
2 vectorizer_tfidf = TfidfVectorizer(min_df=10, ngram_range=(1,4), max_features=5000)
3 vectorizer_tfidf.fit(x_train['preprocessed_essays'].values)# fit has to apply only on
4
5 # we use fitted CountVectorizer to convert the text to vector
6 x_train_tfidf_essays = vectorizer.transform(x_train['preprocessed_essays'].values)
7 x_cv_tfidf_essays = vectorizer.transform(x_cv['preprocessed_essays'].values)
8 x_test_tfidf_essays = vectorizer.transform(x_test['preprocessed_essays'].values)
9
10 print("Shape of matrix after one hot encodig ",x_train_tfidf_essays.shape, y_train.sha
11 print("Shape of matrix after one hot encodig ",x_cv_tfidf_essays.shape)
12 print("Shape of matrix after one hot encodig ",x_test_tfidf_essays.shape)

```

Shape of matrix after one hot encodig (24500, 2142) (24500,)

Shape of matrix after one hot encodig (10500, 2142)

Shape of matrix after one hot encodig (15000, 2142)

In [84]:

```

1 #TFIDF Vectorizer on `project_title`
2
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 vectorizer_tfidf = TfidfVectorizer(min_df=10, ngram_range=(1,4), max_features=5000)
5 vectorizer_tfidf.fit(x_train['preprocessed_project_title'].values)# fit has to apply on
6
7 # we use fitted CountVectorizer to convert the text to vector
8 x_train_tfidf_title = vectorizer.transform(x_train['preprocessed_project_title'].value
9 x_cv_tfidf_title = vectorizer.transform(x_cv['preprocessed_project_title'].values)
10 x_test_tfidf_title = vectorizer.transform(x_test['preprocessed_project_title'].values)
11
12 print("Shape of matrix after one hot encodig ",x_train_tfidf_title.shape)
13 print("Shape of matrix after one hot encodig ",x_cv_tfidf_title.shape)
14 print("Shape of matrix after one hot encodig ",x_test_tfidf_title.shape)
15
16

```

Shape of matrix after one hot encodig (24500, 2142)

Shape of matrix after one hot encodig (10500, 2142)

Shape of matrix after one hot encodig (15000, 2142)

**merge all sparse data, SET 2**



In [85]:

```
1 # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
2 from scipy.sparse import hstack
3 # with the same hstack function we are concatenating a sparse matrix and a dense matrix
4 x_train_tfidf = hstack((x_train_ohe, x_train_tfidf_essays, x_train_tfidf_title)).tocsr()
5 x_cv_tfidf = hstack((x_cv_ohe, x_cv_tfidf_essays, x_cv_tfidf_title)).tocsr()
6 x_test_tfidf = hstack((x_test_ohe, x_test_tfidf_essays, x_test_tfidf_title)).tocsr()
7
8 print(x_train_tfidf.shape)
9 print(x_cv_tfidf.shape)
10 print(x_test_tfidf.shape)
```

(24500, 4384)

(10500, 4384)

(15000, 4384)

In [86]:

```
1 type(x_train_tfidf)
```

Out[86]:

scipy.sparse.csr.csr\_matrix

## Hyperparameter tuning by AUC plot for cv and train dataset, SET 2



You need to plot the performance of model both on train data and cross validation data for each hyper parameter, as shown in the figure

Once you find the best hyper parameter, you need to train your model-M using the best hyper-param. Now, find the AUC on test data and plot the ROC curve on both train and test using model-M.


Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels of test data points

In [124]:

```

1 # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
2 from sklearn.metrics import roc_curve, auc
3
4 neigh = KNeighborsClassifier(n_neighbors=opt_k_tfidf, n_jobs=-1)
5 neigh.fit(x_train_tfidf, y_train)
6 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
7 # not the predicted outputs
8
9 y_train_pred = batch_predict(neigh, x_train_tfidf)
10 y_test_pred = batch_predict(neigh, x_test_tfidf)
11
12 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
13 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
14 tain_auc_tfidf=auc(train_fpr, train_tpr)
15 test_auc_tfidf=auc(test_fpr, test_tpr)
16
17 plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
18 plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
19 plt.legend()
20 plt.xlabel("K: hyperparameter")
21 plt.ylabel("AUC")
22 plt.title("ERROR PLOTS")
23 plt.grid()
24 plt.show()

```



## CONFUSION MATRIX

In [125]:

```
1 #CONFUSION MATRIX
2 def predict(proba, threshold, fpr, tpr):
3
4 t = threshold[np.argmax(fpr*(1-tpr))]
5
6 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
7
8 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.ro
9 predictions = []
10 for i in proba:
11 if i>=t:
12 predictions.append(1)
13 else:
14 predictions.append(0)
15 return predictions
16
17
18
19
```

In [126]:

```

1 print("="*100)
2 print("TRAIN confusion matrix")
3 print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, train_
4
5
6 conf_matr_df_train_1 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pred, tr
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```

TRAIN confusion matrix

the maximum value of  $tpr \cdot (1 - fpr)$  0.3852374534828573 for threshold 0.811


```
[[2471 1231]
```

```
 [9044 11754]]
```

the maximum value of  $tpr \cdot (1 - fpr)$  0.3852374534828573 for threshold 0.811

Out[126]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea218deb00>



In [127]:

```

1 print("="*100)
2 print("Test confusion matrix")
3 print(confusion_matrix(y_test, predict(y_test_pred, te_thresholds, test_fpr, test_tpr))
4
5
6 conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred, te_th
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g')

```


```

=====
=====
Test confusion matrix
the maximum value of tpr*(1-fpr) 0.3404525847312564 for threshold 0.811
[[1415 885]
 [5672 7028]]
the maximum value of tpr*(1-fpr) 0.3404525847312564 for threshold 0.811

```

Out[127]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea211326d8>



## 2.4.3 Applying KNN brute force on AVG W2V, SET 3

vectorize using AVG W2V, SET 3

In [ ]:

```

1 '''
2 # Reading glove vectors in python: https://stackoverflow.com/a/38230349/4084039
3 def loadGloveModel(gloveFile):
4 print ("Loading Glove Model")
5 f = open(gloveFile,'r', encoding="utf8")
6 model = {}
7 for line in tqdm(f):
8 splitLine = line.split()
9 word = splitLine[0]
10 embedding = np.array([float(val) for val in splitLine[1:]])
11 model[word] = embedding
12 print ("Done.",len(model)," words loaded!")
13 return model
14 model = loadGloveModel('glove.42B.300d.txt')
15
16 # =====
17 Output:
18
19 Loading Glove Model
20 1917495it [06:32, 4879.69it/s]
21 Done. 1917495 words loaded!
22
23 # =====
24
25 words = []
26 for i in preprocod_texts:
27 words.extend(i.split(' '))
28
29 for i in preprocod_titles:
30 words.extend(i.split(' '))
31 print("all the words in the coupus", len(words))
32 words = set(words)
33 print("the unique words in the coupus", len(words))
34
35 inter_words = set(model.keys()).intersection(words)
36 print("The number of words that are present in both glove vectors and our coupus", \
37 len(inter_words), "(" ,np.round(len(inter_words)/len(words)*100,3), "%")
38
39 words_courpus = {}
40 words_glove = set(model.keys())
41 for i in words:
42 if i in words_glove:
43 words_courpus[i] = model[i]
44 print("word 2 vec length", len(words_courpus))
45
46
47 # stronging variables into pickle files python: http://www.jessicayung.com/how-to-use-p
48
49 import pickle
50 with open('glove_vectors', 'wb') as f:
51 pickle.dump(words_courpus, f)
52
53
54 '''

```

In [ ]:

```
1 # Please write all the code with proper documentation
```

In [88]:

```
1 # stronging variables into pickle files python: http://www.jessicayung.com/how-to-use-p
2 # make sure you have the glove_vectors file
3 with open('glove_vectors', 'rb') as f:
4 model = pickle.load(f)
5 glove_words = set(model.keys())
6
```



In [90]:

```

1 # -----average Word2Vec on CV
2 # compute average word2vec for each review.
3 avg_w2v_vectors_essays_cv = []; # the avg-w2v for each sentence/review is stored in th
4 for sentence in tqdm(x_cv['preprocessed_essays']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 cnt_words = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if word in glove_words:
9 vector += model[word]
10 cnt_words += 1
11 if cnt_words != 0:
12 vector /= cnt_words
13 avg_w2v_vectors_essays_cv.append(vector)
14
15 print(len(avg_w2v_vectors_essays_cv))
16 print(len(avg_w2v_vectors_essays_cv[0]))

```

```
100%|██|
██████████ 10500/10500 [00:02<00:00, 4718.00it/s]
```

10500  
300

In [91]:

```

1 # -----average Word2Vec on test
2 # compute average word2vec for each review.
3 avg_w2v_vectors_essays_test = []; # the avg-w2v for each sentence/review is stored in
4 for sentence in tqdm(x_test['preprocessed_essays']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 cnt_words = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if word in glove_words:
9 vector += model[word]
10 cnt_words += 1
11 if cnt_words != 0:
12 vector /= cnt_words
13 avg_w2v_vectors_essays_test.append(vector)
14
15 print(len(avg_w2v_vectors_essays_test))
16 print(len(avg_w2v_vectors_essays_test[0]))

```

```
100%|███|
██████████ 15000/15000 [00:03<00:00, 4661.69it/s]
```

15000  
300

In [92]:

```

1 # Using Pretrained Models: AVG W2V on `project_title`
2 # _____
3
4 # -----average Word2Vec on train
5 # compute average word2vec for each review.
6 avg_w2v_vectors_project_title_train = []; # the avg-w2v for each sentence/review is stored
7 for sentence in tqdm(x_train['preprocessed_project_title']): # for each review/sentence
8 vector = np.zeros(300) # as word vectors are of zero length
9 cnt_words = 0; # num of words with a valid vector in the sentence/review
10 for word in sentence.split(): # for each word in a review/sentence
11 if word in glove_words:
12 vector += model[word]
13 cnt_words += 1
14 if cnt_words != 0:
15 vector /= cnt_words
16 avg_w2v_vectors_project_title_train.append(vector)
17
18 print(len(avg_w2v_vectors_project_title_train))
19 print(len(avg_w2v_vectors_project_title_train[0]))
20
21
22
23

```

```
100%|██████████| 24500/24500 [00:00<00:00, 87814.06it/s]
```

24500  
300

In [93]:

```

1 # -----average Word2Vec on cv
2 # compute average word2vec for each review.
3 avg_w2v_vectors_project_title_cv = []; # the avg-w2v for each sentence/review is stored
4 for sentence in tqdm(x_cv['preprocessed_project_title']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 cnt_words = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if word in glove_words:
9 vector += model[word]
10 cnt_words += 1
11 if cnt_words != 0:
12 vector /= cnt_words
13 avg_w2v_vectors_project_title_cv.append(vector)
14
15 print(len(avg_w2v_vectors_project_title_cv))
16 print(len(avg_w2v_vectors_project_title_cv[0]))

```

```
100% |██|
█ | 10500/10500 [00:00<00:00, 84210.73it/s]
```

10500  
300




In [97]:

```

1 train_auc = []
2 cv_auc = []
3 K = [1, 5, 10, 15, 20, 30, 40, 50, 60, 70, 80, 90, 99]
4 for i in tqdm(K):
5 neigh = KNeighborsClassifier(n_neighbors=i, n_jobs=-1)
6 neigh.fit(x_train_AVGW2V, y_train)
7
8
9
10 y_train_pred = batch_predict(neigh, x_train_AVGW2V)
11 y_cv_pred = batch_predict(neigh, x_cv_AVGW2V)
12
13 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates
14 # not the predicted outputs
15 train_auc.append(roc_auc_score(y_train, y_train_pred))
16 cv_auc.append(roc_auc_score(y_cv, y_cv_pred))
17
18 plt.plot(K, train_auc, label='Train AUC')
19 plt.plot(K, cv_auc, label='CV AUC')
20
21 plt.scatter(K, train_auc, label='Train AUC points')
22 plt.scatter(K, cv_auc, label='CV AUC points')
23
24 plt.legend()
25 plt.xlabel("K: hyperparameter")
26 plt.ylabel("AUC")
27 plt.title("ERROR PLOTS")
28 plt.grid()
29 plt.show()

```

```
100%|███████████| 13/13 [2:10:49<00:00, 609.83s/it]
```



In [118]:

1 opt k AVGW2V=90

In [ ]:

1

## Apply best hyperparameter on test dataset, SET 3


You need to plot the performance of model both on train data and cross validation data for each hyperparameter, as shown in the figure

Once you find the best hyper parameter, you need to train your model-M using the best hyper-param. Now, find the AUC on test data and plot the ROC curve on both train and test using model-M.

Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels of test data points

In [119]:

```
1 # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
2 from sklearn.metrics import roc_curve, auc
3
4 neigh = KNeighborsClassifier(n_neighbors=opt_k_AVGW2V, n_jobs=-1)
5 neigh.fit(x_train_AVGW2V, y_train)
6 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the positive class
7 # not the predicted outputs
8
9 y_train_pred = batch_predict(neigh, x_train_AVGW2V)
10 y_test_pred = batch_predict(neigh, x_test_AVGW2V)
11
12 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
13 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
14 train_auc_AVGW2V=auc(train_fpr, train_tpr)
15 test_auc_AVGW2V=auc(test_fpr, test_tpr)
16
17 plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
18 plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
19 plt.legend()
20 plt.xlabel("K: hyperparameter")
21 plt.ylabel("AUC")
22 plt.title("ERROR PLOTS")
23 plt.grid()
24 plt.show()
```



## CONFUSION MATRIX, SET 3

In [120]:

```

1 #CONFUSION MATRIX
2 def predict(proba, threshold, fpr, tpr):
3
4 t = threshold[np.argmax(fpr*(1-tpr))]
5
6 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
7
8 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.ro
9 predictions = []
10 for i in proba:
11 if i>=t:
12 predictions.append(1)
13 else:
14 predictions.append(0)
15 return predictions
16
17
18

```

In [121]:

```

1 print("="*100)
2 print("TRAIN confusion matrix")
3 print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, train_
4
5
6 conf_matr_df_train_1 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pred, tr
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```


```

TRAIN confusion matrix
the maximum value of tpr*(1-fpr) 0.386100440609835 for threshold 0.856
[[2237 1465]
 [7509 13289]]
the maximum value of tpr*(1-fpr) 0.386100440609835 for threshold 0.856

```

Out[121]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1ea218f6cf8>
```



In [122]:

```

1
2 print("="*100)
3 print("Test confusion matrix")
4 print(confusion_matrix(y_test, predict(y_test_pred, te_thresholds, test_fpr, test_tpr))
5
6
7 conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred, te_th
8
9 sns.set(font_scale=1.4)#for label size
10 sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g')

```


```

=====
=====
Test confusion matrix
the maximum value of tpr*(1-fpr) 0.3386686751112633 for threshold 0.856
[[1212 1088]
 [4660 8040]]
the maximum value of tpr*(1-fpr) 0.3386686751112633 for threshold 0.856

```

Out[122]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea2191c470>



## 2.4.4 Applying KNN brute force on TFIDF W2V, SET 4

### Vectorize using TFIDF W2V, SET 4

In [ ]:

```
1 # Please write all the code with proper documentation
```

In [143]:

```

1 # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
2 tfidf_model = TfidfVectorizer()
3 tfidf_model.fit(x_train['preprocessed_essays'])
4 # we are converting a dictionary with word as a key, and the idf as a value
5 dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
6 tfidf_words = set(tfidf_model.get_feature_names())

```

In [144]:

```

1 #Using Pretrained Models: TFIDFW weighted W2V on `essay`
2 # _____
3
4 # average Word2Vec---train
5 # compute average word2vec for each review.
6 tfidf_w2v_vectors_essays_train = []; # the avg-w2v for each sentence/review is stored
7 for sentence in tqdm(x_train['preprocessed_essays']): # for each review/sentence
8 vector = np.zeros(300) # as word vectors are of zero length
9 tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
10 for word in sentence.split(): # for each word in a review/sentence
11 if (word in glove_words) and (word in tfidf_words):
12 vec = model[word] # getting the vector for each word
13 # here we are multiplying idf value(dictionary[word]) and the tf value((sen
14 tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # g
15 vector += (vec * tf_idf) # calculating tfidf weighted w2v
16 tf_idf_weight += tf_idf
17 if tf_idf_weight != 0:
18 vector /= tf_idf_weight
19 tfidf_w2v_vectors_essays_train.append(vector)
20
21 print(len(tfidf_w2v_vectors_essays_train))
22 print(len(tfidf_w2v_vectors_essays_train[0]))
23
24
25
26
27
28

```



In [145]:

```

1 # average Word2Vec---cv
2 # compute average word2vec for each review.
3 tfidf_w2v_vectors_essays_cv = []; # the avg-w2v for each sentence/review is stored in
4 for sentence in tqdm(x_cv['preprocessed_essays']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if (word in glove_words) and (word in tfidf_words):
9 vec = model[word] # getting the vector for each word
10 # here we are multiplying idf value(dictionary[word]) and the tf value((sen
11 tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # g
12 vector += (vec * tf_idf) # calculating tfidf weighted w2v
13 tf_idf_weight += tf_idf
14 if tf_idf_weight != 0:
15 vector /= tf_idf_weight
16 tfidf_w2v_vectors_essays_cv.append(vector)
17
18 print(len(tfidf_w2v_vectors_essays_cv))
19 print(len(tfidf_w2v_vectors_essays_cv[0]))

```

```
100%|██████████| 10500/10500 [00:18<00:00, 582.98it/s]
```

10500  
300

In [146]:

```

1 # average Word2Vec---test
2 # compute average word2vec for each review.
3 tfidf_w2v_vectors_essays_test = []; # the avg-w2v for each sentence/review is stored in
4 for sentence in tqdm(x_test['preprocessed_essays']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if (word in glove_words) and (word in tfidf_words):
9 vec = model[word] # getting the vector for each word
10 # here we are multiplying idf value(dictionary[word]) and the tf value((sentence.count(word)/len(sentence.split()))) # getting the tfidf value
11 tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tfidf value
12 vector += (vec * tf_idf) # calculating tfidf weighted w2v
13 tf_idf_weight += tf_idf
14 if tf_idf_weight != 0:
15 vector /= tf_idf_weight
16 tfidf_w2v_vectors_essays_test.append(vector)
17
18 print(len(tfidf_w2v_vectors_essays_test))
19 print(len(tfidf_w2v_vectors_essays_test[0]))

```

```
100%|███|
██████| 15000/15000 [00:25<00:00, 592.62it/s]
```

15000  
300



In [149]:

```

1 # average Word2Vec--cv
2 # compute average word2vec for each review.
3 tfidf_w2v_vectors_project_title_cv = []; # the avg-w2v for each sentence/review is stored
4 for sentence in tqdm(x_cv['preprocessed_project_title']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if (word in glove_words) and (word in tfidf_words):
9 vec = model[word] # getting the vector for each word
10 # here we are multiplying idf value(dictionary[word]) and the tf value((sentence.count(word)/len(sentence.split()))) # getting the tfidf value
11 tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tfidf value
12 vector += (vec * tf_idf) # calculating tfidf weighted w2v
13 tf_idf_weight += tf_idf
14 if tf_idf_weight != 0:
15 vector /= tf_idf_weight
16 tfidf_w2v_vectors_project_title_cv.append(vector)
17
18 print(len(tfidf_w2v_vectors_project_title_cv))
19 print(len(tfidf_w2v_vectors_project_title_cv[0]))

```

```
100%|██████████| 10500/10500 [00:00<00:00, 40179.61it/s]
```

10500  
300

In [150]:

```

1 # average Word2Vec--test
2 # compute average word2vec for each review.
3 tfidf_w2v_vectors_project_title_test = []; # the avg-w2v for each sentence/review is s
4 for sentence in tqdm(x_test['preprocessed_project_title']): # for each review/sentence
5 vector = np.zeros(300) # as word vectors are of zero length
6 tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
7 for word in sentence.split(): # for each word in a review/sentence
8 if (word in glove_words) and (word in tfidf_words):
9 vec = model[word] # getting the vector for each word
10 # here we are multiplying idf value(dictionary[word]) and the tf value((sen
11 tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # g
12 vector += (vec * tf_idf) # calculating tfidf weighted w2v
13 tf_idf_weight += tf_idf
14 if tf_idf_weight != 0:
15 vector /= tf_idf_weight
16 tfidf_w2v_vectors_project_title_test.append(vector)
17
18 print(len(tfidf_w2v_vectors_project_title_test))
19 print(len(tfidf_w2v_vectors_project_title_test[0]))

```

```
100%|███
██████████ | 15000/15000 [00:00<00:00, 41544.08it/s]
```

15000  
300

## merge all aparse data, SET 4

In [151]:

```
1 # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
2 from scipy.sparse import hstack
3 # with the same hstack function we are concatenating a sparse matrix and a dense matrix
4 x_train_TFIDFW2V = hstack((x_train_ohe, tfidf_w2v_vectors_essays_train, tfidf_w2v_vectors_projects_train))
5 x_cv_TFIDFW2V = hstack((x_cv_ohe, tfidf_w2v_vectors_essays_cv, tfidf_w2v_vectors_projects_cv))
6 x_test_TFIDFW2V = hstack((x_test_ohe, tfidf_w2v_vectors_essays_test, tfidf_w2v_vectors_projects_test))
7
8 print(x_train_TFIDFW2V.shape)
9 print(x_cv_TFIDFW2V.shape)
10 print(x_test_TFIDFW2V.shape)
```

(24500, 700)

(10500, 700)

(15000, 700)

In [152]:

```
1 type(x_train_TFIDFW2V)
```

Out[152]:

scipy.sparse.csr.csr\_matrix

**Hyperparameter tuning by AUC plot for cv and train dataset, SET 4**




You need to plot the performance of model both on train data and cross validation data for each hyper parameter, as shown in the figure

Once you find the best hyper parameter, you need to train your model-M using the best hyper-param. Now, find the AUC on test data and plot the ROC curve on both train and test using model-M.

Along with plotting ROC curve, you need to print the confusion matrix with predicted and original labels of test data points

In [155]:

```
1 # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
2 from sklearn.metrics import roc_curve, auc
3
4 neigh = KNeighborsClassifier(n_neighbors=opt_k_TFIDFW2V, n_jobs=-1)
5 neigh.fit(x_train_TFIDFW2V, y_train)
6 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of
7 # not the predicted outputs
8
9 y_train_pred = batch_predict(neigh, x_train_TFIDFW2V)
10 y_test_pred = batch_predict(neigh, x_test_TFIDFW2V)
11
12 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
13 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
14 tain_auc_TFIDFW2V=auc(train_fpr, train_tpr)
15 test_auc_TFIDFW2V=auc(test_fpr, test_tpr)
16
17 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
18 plt.plot(test_fpr, test_tpr, label="train AUC =" +str(auc(test_fpr, test_tpr)))
19 plt.plot([0,1],[0,1], 'r--')
20 plt.legend()
21 plt.xlabel("K: hyperparameter")
22 plt.ylabel("AUC")
23 plt.title("ERROR PLOTS")
24 plt.grid()
25 plt.show()
```



## CONFUSION MATRIX, SET 4

In [156]:

```

1 def predict(proba, threshold, fpr, tpr):
2
3 t = threshold[np.argmax(fpr*(1-tpr))]
4
5 # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
6
7 print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.rou
8 predictions = []
9 for i in proba:
10 if i>=t:
11 predictions.append(1)
12 else:
13 predictions.append(0)
14 return predictions
15
16
17
18

```

In [157]:

```

1 print("="*100)
2 print("TRAIN confusion matrix")
3 print(confusion_matrix(y_train, predict(y_train_pred, tr_thresholds, train_fpr, train_
4
5
6 conf_matr_df_train_1 = pd.DataFrame(confusion_matrix(y_train, predict(y_train_pred, tr
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_train_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```


```

TRAIN confusion matrix
the maximum value of tpr*(1-fpr) 0.39289929074653884 for threshold 0.85
[[2217 1485]
 [7153 13645]]
the maximum value of tpr*(1-fpr) 0.39289929074653884 for threshold 0.85

```

Out[157]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x1ea2ad3bd30>
```



In [158]:

```

1 print("="*100)
2 print("Test confusion matrix")
3 print(confusion_matrix(y_test, predict(y_test_pred, te_thresholds, test_fpr, test_tpr))
4
5
6 conf_matr_df_test_1 = pd.DataFrame(confusion_matrix(y_test, predict(y_test_pred, te_th
7
8 sns.set(font_scale=1.4)#for label size
9 sns.heatmap(conf_matr_df_test_1, annot=True,annot_kws={"size": 16}, fmt='g')

```

```

=====
=====

```

Test confusion matrix

the maximum value of  $tpr \cdot (1 - fpr)$  0.35439233139335846 for threshold 0.85


```
[[1248 1052]
```

```
[4539 8161]]
```

the maximum value of  $tpr \cdot (1 - fpr)$  0.35439233139335846 for threshold 0.85

Out[158]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x1ea2042c0b8>



## Conclusion

In [ ]:

```
1 # Please compare all your models using Prettytable library
```



In [161]:

```

1 from prettytable import PrettyTable
2 x = PrettyTable()
3 x.field_names = ["Vectorizer", "Model", "Hyper parameter k", "AUC_train", "AUC_test"]
4 x.add_row(["BOW", "Brute", opt_k_bow, tain_auc_bow, test_auc_bow])
5 x.add_row(["BOW_top2000", "Brute", opt_k_bow_2000, tain_auc_bow_2000, test_auc_bow_2000])
6 x.add_row(["TFIDF", "Brute", opt_k_tfidf, tain_auc_tfidf, test_auc_tfidf])
7 x.add_row(["AVG W2V", "Brute", opt_k_AVGW2V, tain_auc_AVGW2V, test_auc_AVGW2V])
8 x.add_row(["TFIDF W2V", "Brute", opt_k_TFIDFW2V, tain_auc_TFIDFW2V, test_auc_TFIDFW2V])
9 print(x)
10 with open('KNN_Result.txt', 'w') as file:
11 file.write(str(x))
12

```

```

+-----+-----+-----+-----+-----+
| Vectorizer | Model | Hyper parameter k | AUC_train | AUC_test |
+-----+-----+-----+-----+-----+
BOW	Brute	90	0.6675832032845697	0.614286340
BOW_top2000	Brute	40	0.6966539075231073	0.514796953
TFIDF	Brute	90	0.6662211395778457	0.666221139
AVG W2V	Brute	90	0.6677880304120587	0.667788030
TFIDF W2V	Brute	80	0.6771984670117213	0.630027986
+-----+-----+-----+-----+-----+

```

**Observation:** (Above analysis done on 50000 dataset due to memory issue)

1) Hyper parameter value k is same i.e 3 for all cases.

2) there is not much difference in AUC value of train and test value and AUC value not much low that shows that model is neither underfit nor overfit.

3) most of error occurring in FN box of confusion where actual value is 1 but predicted 0.

In [ ]:

1

In [ ]:

1

In [ ]:

1

